
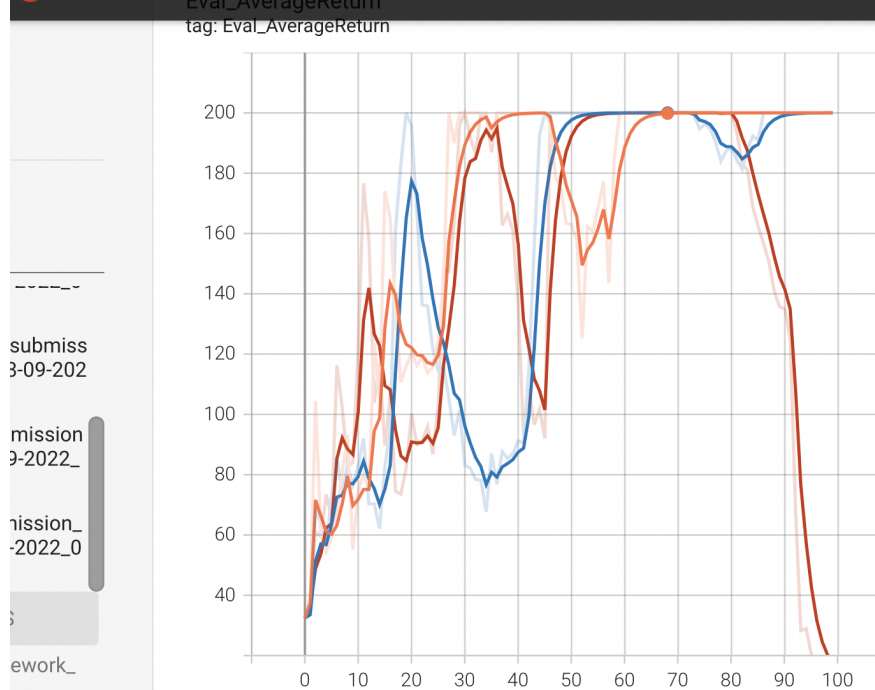


Experiment 1

q1_sb_no_rtg_dsa_submission_CartPole-v0_28-09-2022_07-26-02	200	200	93	Wed Sep 28, 00:27:42	1m 35s
q1_sb_rtg_dsa_submission_CartPole-v0_28-09-2022_07-27-57	178.5	175.3	93	Wed Sep 28, 00:29:40	1m 40s
q1_sb_rtg_na_submission_CartPole-v0_28-09-2022_07-30-12	58.29	52	93	Wed Sep 28, 00:31:53	1m 39s



q1_lb_no_rtg_dsa_submission_CartPole-v0_28-09-2022_07-32-41	199.8	200	68	Wed Sep 28, 00:36:54	4m 9s
q1_lb_rtg_dsa_submission_CartPole-v0_28-09-2022_07-40-47	200	200	68	Wed Sep 28, 00:44:59	4m 8s
q1_lb_rtg_na_submission_CartPole-v0_28-09-2022_07-48-48	200	200	68	Wed Sep 28, 00:52:57	4m 4s




Experiment 2

for best performance:

$b = 500$

$r = 1e-2$

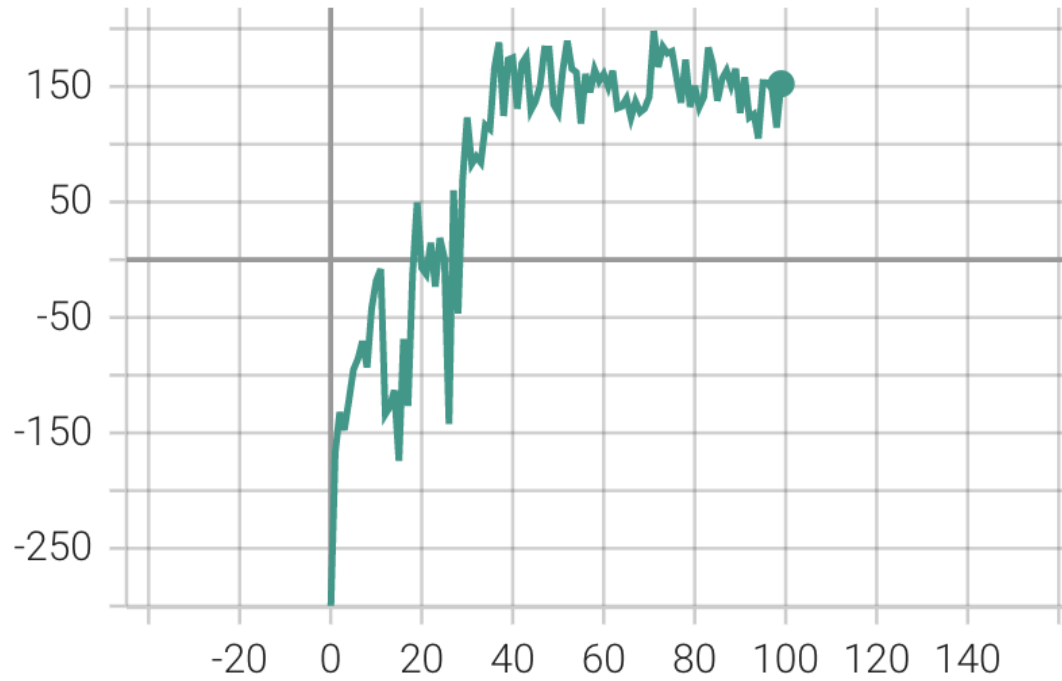
Eval_AverageReturn
tag: Eval_AverageReturn



Commands see "submission_commands.txt"


Experiment 3

Eval_AverageReturn
tag: Eval_AverageReturn



Experiment 4

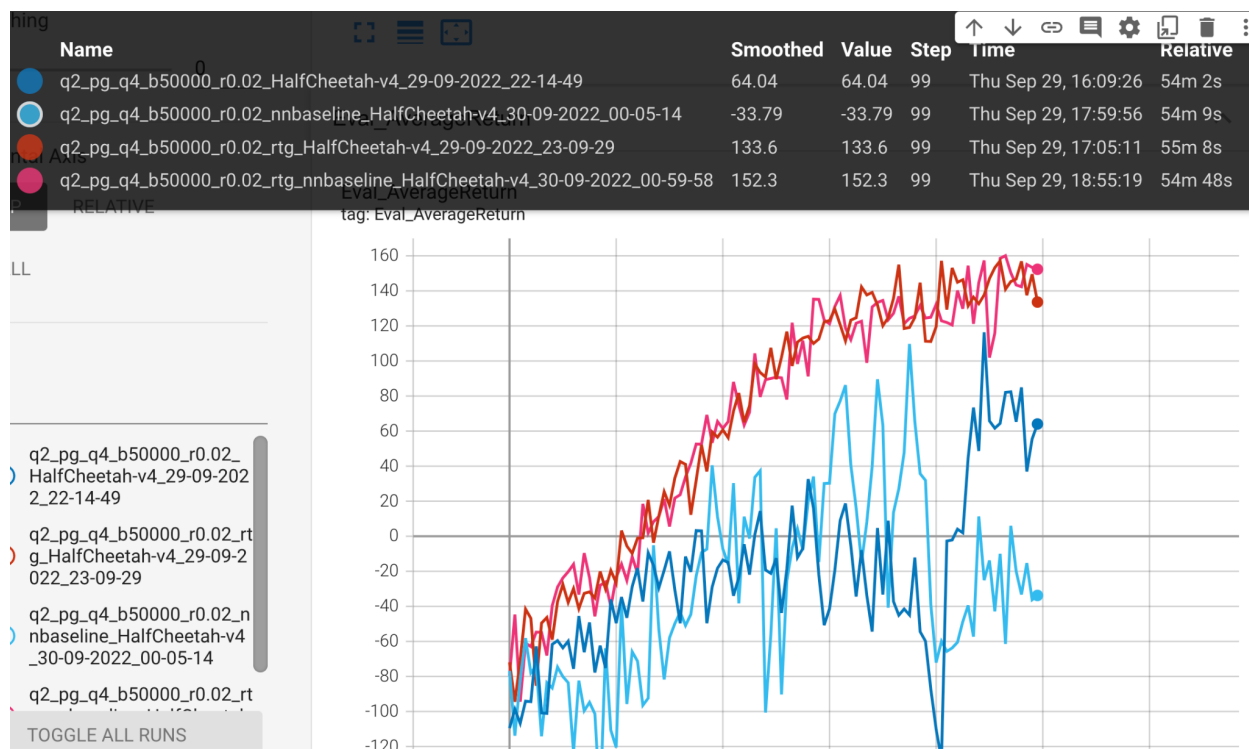
Eval_AverageReturn
tag: Eval_AverageReturn



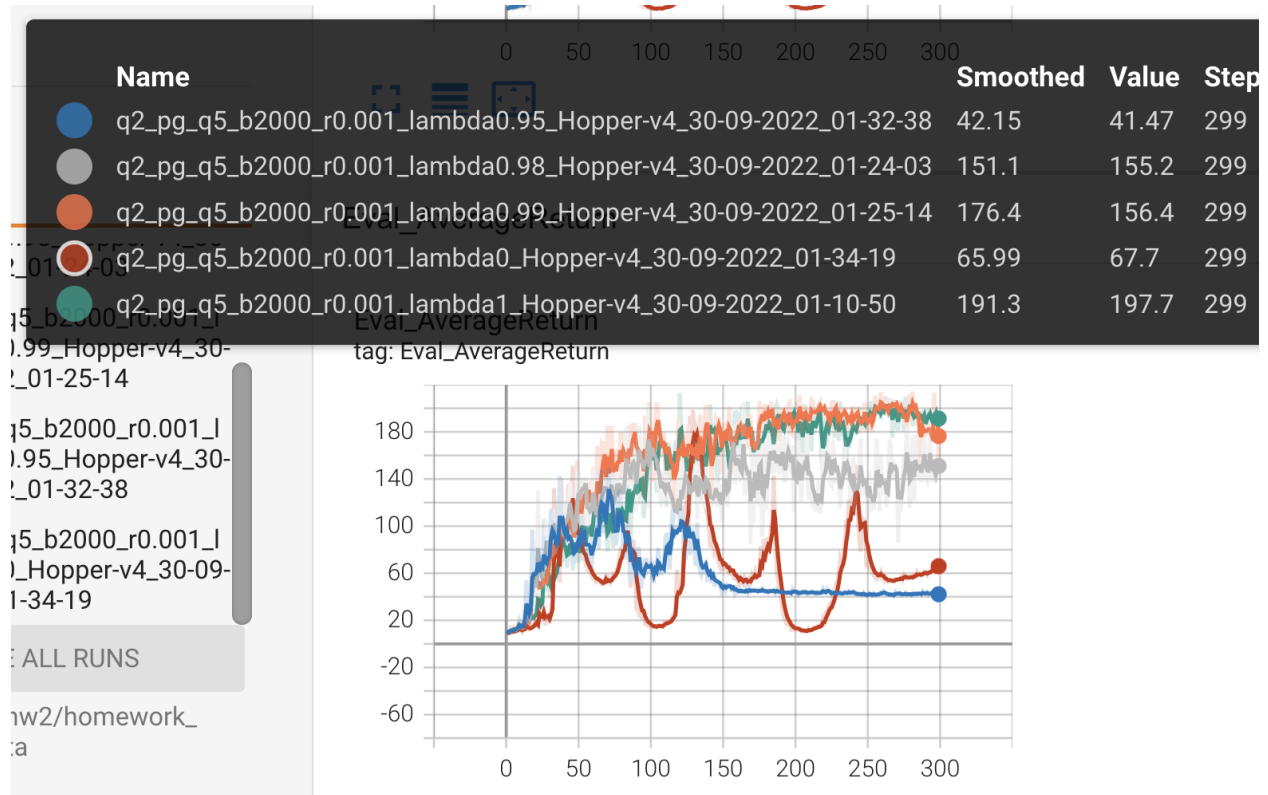
Legend:

Name	Smoothed	Value	Step		
q2_pg_q4_search_b10000_lr0.005_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_21-47-59	2.475	2.475	99	Thu Sep 29, 14:59:41	11m 34s
q2_pg_q4_search_b10000_lr0.01_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_15-34-34	73.74	73.74	99	Thu Sep 29, 08:45:39	10m 57s
q2_pg_q4_search_b10000_lr0.02_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_22-00-38	27.37	27.37	99	Thu Sep 29, 15:12:47	12m 1s
q2_pg_q4_search_b30000_lr0.005_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_20-42-46	22.69	22.69	99	Thu Sep 29, 14:16:52	33m 44s
q2_pg_q4_search_b30000_lr0.01_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_06-50-13	130.6	130.6	99	Thu Sep 29, 00:39:26	48m 39s
q2_pg_q4_search_b30000_lr0.02_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_20-08-07	132.5	132.5	99	Thu Sep 29, 13:41:40	33m 13s
q2_pg_q4_search_b50000_lr0.005_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_17-03-34	63.64	63.64	99	Thu Sep 29, 10:58:53	54m 46s
q2_pg_q4_search_b50000_lr0.01_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_16-04-46	110.3	110.3	99	Thu Sep 29, 09:59:29	54m 8s
q2_pg_q4_search_b50000_lr0.02_rtg_nnbaseline_HalfCheetah-v4_29-09-2022_19-05-29	152.3	152.3	99	Thu Sep 29, 13:00:49	54m 46s

The results shows that among {0.005, 0.01, 0.02}, the larger the learning rate the better the result. Larger batch size seems to also be better but the effect in this experiment is limited comparing to the learning rate.



Experiment 5



Higher lambda values stabilize the training process.

#@title runtime arguments

class Args:

```
def __getitem__(self, key):  
    return getattr(self, key)
```

```
def __setitem__(self, key, val):  
    setattr(self, key, val)
```

```
def __contains__(self, key):  
    return hasattr(self, key)
```

```
env_name = 'CartPole-v0' #@param  
exp_name = 'q1_sb_no_rtg_dsa_submission' #@param
```

```
#@markdown main parameters of interest  
n_iter = 100 #@param {type: "integer"}
```

```
## PDF will tell you how to set ep_len  
## and discount for each environment  
ep_len = 200 #@param {type: "integer"}  
discount = 1 #@param {type: "number"}
```

```
reward_to_go = False #@param {type: "boolean"}  
nn_baseline = False #@param {type: "boolean"}  
gae_lambda = None #@param {type: "float"}  
dont_standardize_advantages = False #@param {type: "boolean"}
```

```
#@markdown batches and steps  
batch_size = 1000 #@param {type: "integer"}  
eval_batch_size = 400 #@param {type: "integer"}
```

```
num_agent_train_steps_per_iter = 1 #@param {type: "integer"}  
learning_rate = 5e-3 #@param {type: "number"}
```

```
#@markdown MLP parameters  
n_layers = 2 #@param {type: "integer"}  
size = 64 #@param {type: "integer"}
```

```
#@markdown system  
save_params = False #@param {type: "boolean"}  
no_gpu = False #@param {type: "boolean"}  
which_gpu = 0 #@param {type: "integer"}  
seed = 1 #@param {type: "integer"}
```

```
action_noise_std = 0 #@param {type: "float"}
```

```

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )
-----
-----
#@title runtime arguments

class Args:

    def __getitem__(self, key):
        return getattr(self, key)

    def __setitem__(self, key, val):
        setattr(self, key, val)

    def __contains__(self, key):
        return hasattr(self, key)

    env_name = 'CartPole-v0' #@param
    exp_name = 'q1_sb_rtg_dsa_submission' #@param

    #@markdown main parameters of interest
    n_iter = 100 #@param {type: "integer"}

    ## PDF will tell you how to set ep_len
    ## and discount for each environment
    ep_len = 200 #@param {type: "integer"}
    discount = 1 #@param {type: "number"}

    reward_to_go = True #@param {type: "boolean"}
    nn_baseline = False #@param {type: "boolean"}
    gae_lambda = None #@param {type: "float"}
    dont_standardize_advantages = False #@param {type: "boolean"}

```



```

#@markdown batches and steps
batch_size = 1000 #@param {type: "integer"}
eval_batch_size = 400 #@param {type: "integer"}

num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
learning_rate = 5e-3 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )
    -----
    -----
#@title runtime arguments

class Args:

    def __getitem__(self, key):
        return getattr(self, key)

    def __setitem__(self, key, val):
        setattr(self, key, val)

```

```

def __contains__(self, key):
    return hasattr(self, key)

env_name = 'CartPole-v0' #@param
exp_name = 'q1_sb_rtg_na_submission' #@param

#@markdown main parameters of interest
n_iter = 100 #@param {type: "integer"}

## PDF will tell you how to set ep_len
## and discount for each environment
ep_len = 200 #@param {type: "integer"}
discount = 1 #@param {type: "number"}

reward_to_go = True #@param {type: "boolean"}
nn_baseline = False #@param {type: "boolean"}
gae_lambda = None #@param {type: "float"}
dont_standardize_advantages = True #@param {type: "boolean"}

#@markdown batches and steps
batch_size = 1000 #@param {type: "integer"}
eval_batch_size = 400 #@param {type: "integer"}

num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
learning_rate = 5e-3 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

```

```

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )

```

```

#@title runtime arguments

```

```

class Args:

```

```

    def __getitem__(self, key):
        return getattr(self, key)

```

```

    def __setitem__(self, key, val):
        setattr(self, key, val)

```

```

    def __contains__(self, key):
        return hasattr(self, key)

```

```

    env_name = 'CartPole-v0' #@param
    exp_name = 'q1_lb_no_rtg_dsa_submission' #@param

```

```

    #@markdown main parameters of interest
    n_iter = 100 #@param {type: "integer"}

```

```

    ## PDF will tell you how to set ep_len
    ## and discount for each environment
    ep_len = 200 #@param {type: "integer"}
    discount = 1 #@param {type: "number"}

```

```

    reward_to_go = False #@param {type: "boolean"}
    nn_baseline = False #@param {type: "boolean"}
    gae_lambda = None #@param {type: "float"}
    dont_standardize_advantages = False #@param {type: "boolean"}

```

```

    #@markdown batches and steps
    batch_size = 5000 #@param {type: "integer"}
    eval_batch_size = 400 #@param {type: "integer"}

```

```

    num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
    learning_rate = 5e-3 #@param {type: "number"}

```

```

    #@markdown MLP parameters
    n_layers = 2 #@param {type: "integer"}
    size = 64 #@param {type: "integer"}

```

```

    #@markdown system

```

```

save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )
-----
#@title runtime arguments

class Args:

    def __getitem__(self, key):
        return getattr(self, key)

    def __setitem__(self, key, val):
        setattr(self, key, val)

    def __contains__(self, key):
        return hasattr(self, key)

env_name = 'CartPole-v0' #@param
exp_name = 'q1_lb_rtg_dsa_submission' #@param

#@markdown main parameters of interest
n_iter = 100 #@param {type: "integer"}

## PDF will tell you how to set ep_len
## and discount for each environment
ep_len = 200 #@param {type: "integer"}

```

```

discount = 1 #@param {type: "number"}

reward_to_go = True #@param {type: "boolean"}
nn_baseline = False #@param {type: "boolean"}
gae_lambda = None #@param {type: "float"}
dont_standardize_advantages = False #@param {type: "boolean"}

#@markdown batches and steps
batch_size = 5000 #@param {type: "integer"}
eval_batch_size = 400 #@param {type: "integer"}

num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
learning_rate = 5e-3 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )

```

```

#@title runtime arguments

class Args:

```

```

def __getitem__(self, key):
    return getattr(self, key)

def __setitem__(self, key, val):
    setattr(self, key, val)

def __contains__(self, key):
    return hasattr(self, key)

env_name = 'CartPole-v0' #@param
exp_name = 'q1_lb_rtg_na_submission' #@param

#@markdown main parameters of interest
n_iter = 100 #@param {type: "integer"}

## PDF will tell you how to set ep_len
## and discount for each environment
ep_len = 200 #@param {type: "integer"}
discount = 1 #@param {type: "number"}

reward_to_go = True #@param {type: "boolean"}
nn_baseline = False #@param {type: "boolean"}
gae_lambda = None #@param {type: "float"}
dont_standardize_advantages = True #@param {type: "boolean"}

#@markdown batches and steps
batch_size = 5000 #@param {type: "integer"}
eval_batch_size = 400 #@param {type: "integer"}

num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
learning_rate = 5e-3 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}

```

```

    scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )
-----
#@title runtime arguments

class Args:

    def __getitem__(self, key):
        return getattr(self, key)

    def __setitem__(self, key, val):
        setattr(self, key, val)

    def __contains__(self, key):
        return hasattr(self, key)

    env_name = 'InvertedPendulum-v4' #@param
    exp_name = 'q2_b500_r1e-2' #@param

    #@markdown main parameters of interest
    n_iter = 100 #@param {type: "integer"}

    ## PDF will tell you how to set ep_len
    ## and discount for each environment
    ep_len = 1000 #@param {type: "integer"}
    discount = 0.9 #@param {type: "number"}

    reward_to_go = True #@param {type: "boolean"}
    nn_baseline = False #@param {type: "boolean"}
    gae_lambda = None #@param {type: "float"}
    dont_standardize_advantages = False #@param {type: "boolean"}

    #@markdown batches and steps
    batch_size = 500 #@param {type: "integer"}
    eval_batch_size = 400 #@param {type: "integer"}

    num_agent_train_steps_per_iter = 1 #@param {type: "integer"}

```

```

learning_rate = 1e-2 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = True #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 0 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )

```