

---

---

-----  
#@title runtime arguments

class Args:

```
def __getitem__(self, key):  
    return getattr(self, key)
```

```
def __setitem__(self, key, val):  
    setattr(self, key, val)
```

```
def __contains__(self, key):  
    return hasattr(self, key)
```

```
env_name = 'CartPole-v0' #@param  
exp_name = 'q1_sb_no_rtg_dsa_submission' #@param
```

```
#@markdown main parameters of interest  
n_iter = 100 #@param {type: "integer"}
```

```
## PDF will tell you how to set ep_len  
## and discount for each environment  
ep_len = 200 #@param {type: "integer"}  
discount = 1 #@param {type: "number"}
```

```
reward_to_go = False #@param {type: "boolean"}  
nn_baseline = False #@param {type: "boolean"}  
gae_lambda = None #@param {type: "float"}  
dont_standardize_advantages = False #@param {type: "boolean"}
```

```
#@markdown batches and steps  
batch_size = 1000 #@param {type: "integer"}  
eval_batch_size = 400 #@param {type: "integer"}
```

```
num_agent_train_steps_per_iter = 1 #@param {type: "integer"}  
learning_rate = 5e-3 #@param {type: "number"}
```

```
#@markdown MLP parameters  
n_layers = 2 #@param {type: "integer"}  
size = 64 #@param {type: "integer"}
```

```
#@markdown system  
save_params = False #@param {type: "boolean"}  
no_gpu = False #@param {type: "boolean"}  
which_gpu = 0 #@param {type: "integer"}  
seed = 1 #@param {type: "integer"}
```

```
action_noise_std = 0 #@param {type: "float"}
```

```

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )
-----
-----
#@title runtime arguments

class Args:

    def __getitem__(self, key):
        return getattr(self, key)

    def __setitem__(self, key, val):
        setattr(self, key, val)

    def __contains__(self, key):
        return hasattr(self, key)

    env_name = 'CartPole-v0' #@param
    exp_name = 'q1_sb_rtg_dsa_submission' #@param

    #@markdown main parameters of interest
    n_iter = 100 #@param {type: "integer"}

    ## PDF will tell you how to set ep_len
    ## and discount for each environment
    ep_len = 200 #@param {type: "integer"}
    discount = 1 #@param {type: "number"}

    reward_to_go = True #@param {type: "boolean"}
    nn_baseline = False #@param {type: "boolean"}
    gae_lambda = None #@param {type: "float"}
    dont_standardize_advantages = False #@param {type: "boolean"}

```

```

#@markdown batches and steps
batch_size = 1000 #@param {type: "integer"}
eval_batch_size = 400 #@param {type: "integer"}

num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
learning_rate = 5e-3 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )
    -----
    -----
#@title runtime arguments

class Args:

    def __getitem__(self, key):
        return getattr(self, key)

    def __setitem__(self, key, val):
        setattr(self, key, val)

```

```

def __contains__(self, key):
    return hasattr(self, key)

env_name = 'CartPole-v0' #@param
exp_name = 'q1_sb_rtg_na_submission' #@param

#@markdown main parameters of interest
n_iter = 100 #@param {type: "integer"}

## PDF will tell you how to set ep_len
## and discount for each environment
ep_len = 200 #@param {type: "integer"}
discount = 1 #@param {type: "number"}

reward_to_go = True #@param {type: "boolean"}
nn_baseline = False #@param {type: "boolean"}
gae_lambda = None #@param {type: "float"}
dont_standardize_advantages = True #@param {type: "boolean"}

#@markdown batches and steps
batch_size = 1000 #@param {type: "integer"}
eval_batch_size = 400 #@param {type: "integer"}

num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
learning_rate = 5e-3 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

```

```

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )

```

---

```

#@title runtime arguments

```

```

class Args:

```

```

    def __getitem__(self, key):
        return getattr(self, key)

```

```

    def __setitem__(self, key, val):
        setattr(self, key, val)

```

```

    def __contains__(self, key):
        return hasattr(self, key)

```

```

    env_name = 'CartPole-v0' #@param
    exp_name = 'q1_lb_no_rtg_dsa_submission' #@param

```

```

    #@markdown main parameters of interest
    n_iter = 100 #@param {type: "integer"}

```

```

    ## PDF will tell you how to set ep_len
    ## and discount for each environment
    ep_len = 200 #@param {type: "integer"}
    discount = 1 #@param {type: "number"}

```

```

    reward_to_go = False #@param {type: "boolean"}
    nn_baseline = False #@param {type: "boolean"}
    gae_lambda = None #@param {type: "float"}
    dont_standardize_advantages = False #@param {type: "boolean"}

```

```

    #@markdown batches and steps
    batch_size = 5000 #@param {type: "integer"}
    eval_batch_size = 400 #@param {type: "integer"}

```

```

    num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
    learning_rate = 5e-3 #@param {type: "number"}

```

```

    #@markdown MLP parameters
    n_layers = 2 #@param {type: "integer"}
    size = 64 #@param {type: "integer"}

```

```

    #@markdown system

```

```

save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )
-----
#@title runtime arguments

class Args:

    def __getitem__(self, key):
        return getattr(self, key)

    def __setitem__(self, key, val):
        setattr(self, key, val)

    def __contains__(self, key):
        return hasattr(self, key)

env_name = 'CartPole-v0' #@param
exp_name = 'q1_lb_rtg_dsa_submission' #@param

#@markdown main parameters of interest
n_iter = 100 #@param {type: "integer"}

## PDF will tell you how to set ep_len
## and discount for each environment
ep_len = 200 #@param {type: "integer"}

```

```

discount = 1 #@param {type: "number"}

reward_to_go = True #@param {type: "boolean"}
nn_baseline = False #@param {type: "boolean"}
gae_lambda = None #@param {type: "float"}
dont_standardize_advantages = False #@param {type: "boolean"}

#@markdown batches and steps
batch_size = 5000 #@param {type: "integer"}
eval_batch_size = 400 #@param {type: "integer"}

num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
learning_rate = 5e-3 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )

```

---

```

#@title runtime arguments

class Args:

```

```

def __getitem__(self, key):
    return getattr(self, key)

def __setitem__(self, key, val):
    setattr(self, key, val)

def __contains__(self, key):
    return hasattr(self, key)

env_name = 'CartPole-v0' #@param
exp_name = 'q1_lb_rtg_na_submission' #@param

#@markdown main parameters of interest
n_iter = 100 #@param {type: "integer"}

## PDF will tell you how to set ep_len
## and discount for each environment
ep_len = 200 #@param {type: "integer"}
discount = 1 #@param {type: "number"}

reward_to_go = True #@param {type: "boolean"}
nn_baseline = False #@param {type: "boolean"}
gae_lambda = None #@param {type: "float"}
dont_standardize_advantages = True #@param {type: "boolean"}

#@markdown batches and steps
batch_size = 5000 #@param {type: "integer"}
eval_batch_size = 400 #@param {type: "integer"}

num_agent_train_steps_per_iter = 1 #@param {type: "integer"}
learning_rate = 5e-3 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = False #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 1 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}

```



```

    scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )
-----
#@title runtime arguments

class Args:

    def __getitem__(self, key):
        return getattr(self, key)

    def __setitem__(self, key, val):
        setattr(self, key, val)

    def __contains__(self, key):
        return hasattr(self, key)

    env_name = 'InvertedPendulum-v4' #@param
    exp_name = 'q2_b500_r1e-2' #@param

    #@markdown main parameters of interest
    n_iter = 100 #@param {type: "integer"}

    ## PDF will tell you how to set ep_len
    ## and discount for each environment
    ep_len = 1000 #@param {type: "integer"}
    discount = 0.9 #@param {type: "number"}

    reward_to_go = True #@param {type: "boolean"}
    nn_baseline = False #@param {type: "boolean"}
    gae_lambda = None #@param {type: "float"}
    dont_standardize_advantages = False #@param {type: "boolean"}

    #@markdown batches and steps
    batch_size = 500 #@param {type: "integer"}
    eval_batch_size = 400 #@param {type: "integer"}

    num_agent_train_steps_per_iter = 1 #@param {type: "integer"}

```

```

learning_rate = 1e-2 #@param {type: "number"}

#@markdown MLP parameters
n_layers = 2 #@param {type: "integer"}
size = 64 #@param {type: "integer"}

#@markdown system
save_params = False #@param {type: "boolean"}
no_gpu = True #@param {type: "boolean"}
which_gpu = 0 #@param {type: "integer"}
seed = 0 #@param {type: "integer"}

action_noise_std = 0 #@param {type: "float"}

#@markdown logging
## default is to not log video so
## that logs are small enough to be
## uploaded to gradscope
video_log_freq = -1#@param {type: "integer"}
scalar_log_freq = 1#@param {type: "integer"}

args = Args()

## ensure compatibility with hw1 code
args['train_batch_size'] = args['batch_size']

if args['video_log_freq'] > 0:
    import warnings
    warnings.warn(
        '''\nLogging videos will make eventfiles too'''
        '''\nlarge for the autograder. Set video_log_freq = -1'''
        '''\nfor the runs you intend to submit.'''
    )

```